

BDS Workshop

Prototyping an Integrated AI Application
From Business Idea to Working POC in 2 Hours

Christian Nielsen
Eskil Olav Andersen, PhD

Aalborg University Business School
Innovation, Knowledge, and Economic Development (IKE)
AI Denmark (AIDK)
eoae@business.aau.dk

November 7, 2025

Introduction

Necessary Steps

Dataset and Contextualization

Technical Components

Evaluation and Execution

Introduction



Build a **working proof-of-concept** application

Your Mission:

- Create a Streamlit app
- Demonstrate clear business value
- Runs locally (no deployment)
- 5-minute stakeholder pitch

Time Available:

- 2 hours of development
- Access to full dataset
- LLMs as co-pilots
- Teacher support

Key Point: Focus: Demonstrate value, not perfect code



What You'll Deliver:

1. **Working Streamlit App** - Runs on your laptop, demonstrates your solution
2. **5-Minute Pitch** - Problem, demo, technical approach, insights
3. **Clear Value Proposition** - Who is this for? What problem does it solve?

What Success Looks Like:

- ✓ App runs without errors
- ✓ Solves a clear business problem
- ✓ You can explain it convincingly
- ✓ Components work together
- ✓ Generates actionable insights

What's NOT Required:

- ✗ Production-ready code
- ✗ Perfect error handling
- ✗ Beautiful UI design
- ✗ Scalable architecture
- ✗ Comprehensive testing

Can you pitch it convincingly in 3-5 minutes?

Proof of Concept

- **Timeline:** Hours to days
- **Goal:** Show it works
- **Quality:** Good enough
- **Scope:** Core idea only
- **Users:** Stakeholders/demo
- **Focus:** *Does this solve the problem?*

If POC succeeds →

Production System

- **Timeline:** Weeks to months
- **Goal:** Ship to users
- **Quality:** Robust, tested
- **Scope:** Full feature set
- **Users:** Real end-users
- **Focus:** *Scalable & maintainable*

Necessary Steps

Start with clarity on WHO and WHY:

Who?

- Apartment seekers?
- Landlords?
- Investors?
- Platform users?

What Problem?

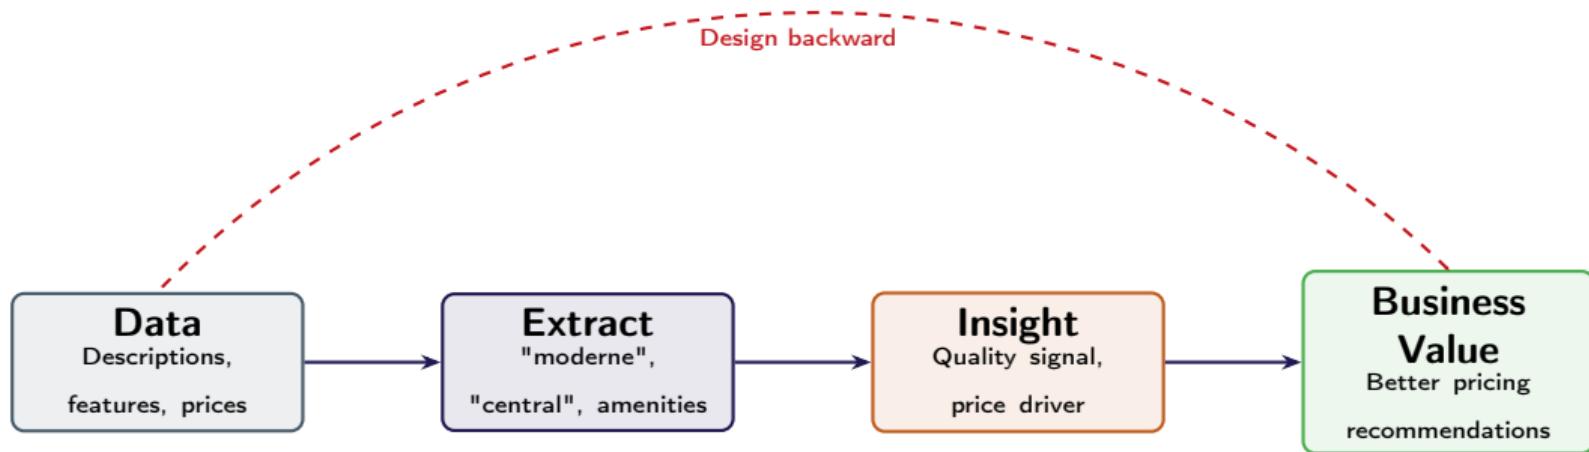
- Finding apartments?
- Pricing decisions?
- Market insights?
- Search quality?

Why Better?

- Faster?
- More accurate?
- New insights?
- Better UX?

Key Point: Clear value proposition = Easier to build & easier to pitch

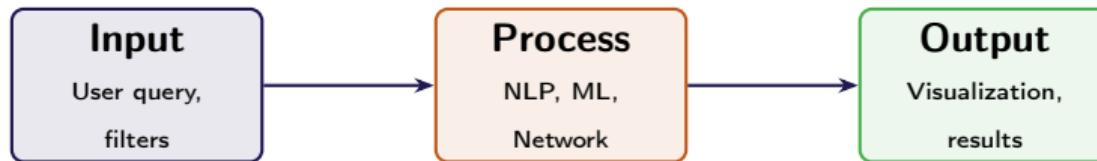
Work backward from the value you want to create:



Example:

- **Goal:** Recommend similar apartments
- **Need:** Similarity measure (cosine, network distance)
- **Requires:** Feature vectors or embeddings from descriptions
- **Therefore:** Extract text features OR use semantic embeddings

Break your app into composable components:



Why Modular?

Benefits:

- Test components independently
- Debug in isolation
- Replace/improve one part
- Easier to explain

Integration Strategy:

- Build components first
- Test with dummy data
- Connect incrementally
- Validate each connection

Important: Components should be composable and testable independently

At 1 Hour: STOP and Evaluate

You should have:

- ✓ Data loads successfully
- ✓ One analysis component works (NLP/M-L/Network)
- ✓ Basic Streamlit app runs

If not, consider:

- Simplify scope
- Use CSV from EDA notebook
- Ask teachers for suggestions
- Adjust ambition level

Dataset and Contextualization

Boligportal Aalborg Dataset

Scale & Structure:

- 1,330 rental apartments
- Scrapped November 2025
- Rich structured fields
- Danish text descriptions
- Complete metadata

Key Fields:

- Rent, size, rooms, floor
- Amenities (elevator, balcony, parking)
- Location (street, postal code)
- Description (avg 1,049 characters)

What You Have:

- boligportal_aalborg.jsonl
- EDA notebook (full exploration)
- CSV export (if preferred)
- Example notebook (clustering, recommendations)

Data Quality:

- ✓ 0% missing critical fields
- ✓ Complete descriptions
- ✓ Clean, ready to use

Different users, different problems, different solutions:

Apartment Seekers:

"Find apartments matching my preferences"

→ Semantic search + recommendations + similarity networks

Landlords:

"Is my apartment priced correctly?"

→ Price prediction + comparable analysis + market positioning

Market Analysts:

"What defines each market segment?"

→ Clustering + feature analysis + segment profiling

Investors:

"Which neighborhoods offer best value?"

→ Location analysis + price-quality ratios + trend identification

Technical Components



Descriptions hold rich details. Start simple, escalate as needed.

- **Basic (Fast & Reliable): Keyword Searching**

e.g., "altan" in `description.lower()`

- ✓ **Pro:** Easy, fast, great for binary features (has `dishwasher`).
- ✗ **Con:** Misses synonyms (`balkon`) and context (`ikke støjende` vs `rolig`).

- **Intermediate (Semantic Power): Embeddings**

e.g., `using sentence-transformers`

- ✓ **Pro:** Captures meaning. Find "quiet modern" from "peaceful renovated".
- ✗ **Con:** Heavier computation than simple keywords.

- **Advanced (Generative AI): Large Language Models (LLMs)**

Use for complex extraction (JSON of amenities), summarization, or classification (e.g., "vibe").

LLM Trade-offs: API vs. Local

API (OpenAI, Google)

- ✓ Extremely powerful
- ✓ Simple to use
- ✗ Network latency
- ✗ Potential costs/limits

Local (Llama, Mistral)

- ✓ No network lag, private
- ✓ Free to run (on own hardware)
- ✗ Requires setup & GPU
- ✗ Models can be less capable



Use ML to answer core business questions directly from your features.

- **Clustering (e.g., K-Means)**

Business Question: "What market segments exist?"

→ Groups similar apartments to create profiles like "Luxury Penthouses" or "Student Studios".

- **Regression (e.g., Random Forest)**

Business Question: "Is this apartment priced correctly?"

→ Predicts rent from features to power a "fair deal" finder or landlord pricing tool.

- **Recommendation (e.g., Cosine Similarity)**

Business Question: "If a user likes this, what else will they like?"

→ Finds the "closest" apartments based on features to improve user discovery.

Networks provide a unique, visual way to understand connections.

Why Use Networks? They excel at showing the *structure* of relationships hidden in the data.

- **Amenity Co-occurrence Network**

Nodes = amenities, Edges = they appear together often.

→ Reveals what defines a "premium" apartment (e.g., "elevator", "balcony", and "parking" form a tight cluster).

- **Apartment Similarity Network**

Nodes = apartments, Edges = high feature similarity.

→ Identifies direct competitors, finds comparables, and reveals market sub-clusters.

- **Location Network**

Nodes = apartments, Edges = on the same street.

→ Analyze hyper-local market dynamics and price variations within a neighborhood.

Note: This can be a powerful differentiator for your pitch, but make sure your core value proposition is working first. It's an extension, not a requirement.

Your UI's job: Make your analysis interactive and your insights understandable.

- **Pattern 1: The Interactive Filter**

The user configures inputs to see a specific slice of the data.

- **Layout:** Controls in a sidebar, results (charts, map, tables) in the main area.
- **Use Case:** Perfect for dashboards and general data exploration.

- **Pattern 2: The "Drill-Down" Explorer**

The user starts with a high-level view and clicks to reveal more detail.

- **Layout:** A single, rich visualization (like a map) is the primary interface. Selecting a point reveals its details.
- **Use Case:** Ideal for geographic analysis or finding specific items in a large dataset.

- **Pattern 3: The "What-If" Simulator**

The user adjusts parameters to see how a model's prediction changes in real-time.

- **Layout:** Dominated by sliders and input boxes that feed directly into your model.
- **Use Case:** Excellent for demonstrating a price predictor or a recommendation engine.

Design Tip: Start by asking: "What question is my user trying to answer?" Then, build the simplest possible interface to help them find that answer.

Good integration means components can be modified or replaced without breaking the whole system.

Principles of Modular Integration:

- **Clear Input/Output Contracts**

Each component should have well-defined inputs and outputs.

- *Example:* NLP component takes raw text → returns feature dictionary
- ✓ You can swap NLP methods (regex → LLM) without touching other code

- **Separation of Concerns**

Keep data processing, analysis, and presentation in separate functions/files.

- *Example:* data.py, analysis.py, app.py
- ✓ Debug ML model without worrying about UI, fix UI without touching ML

- **Test Components Independently**

Build and validate each component with dummy/sample data first.

- *Example:* Test clustering on 50 apartments before running on full dataset
- ✓ Catch bugs early, iterate faster, integration is easier

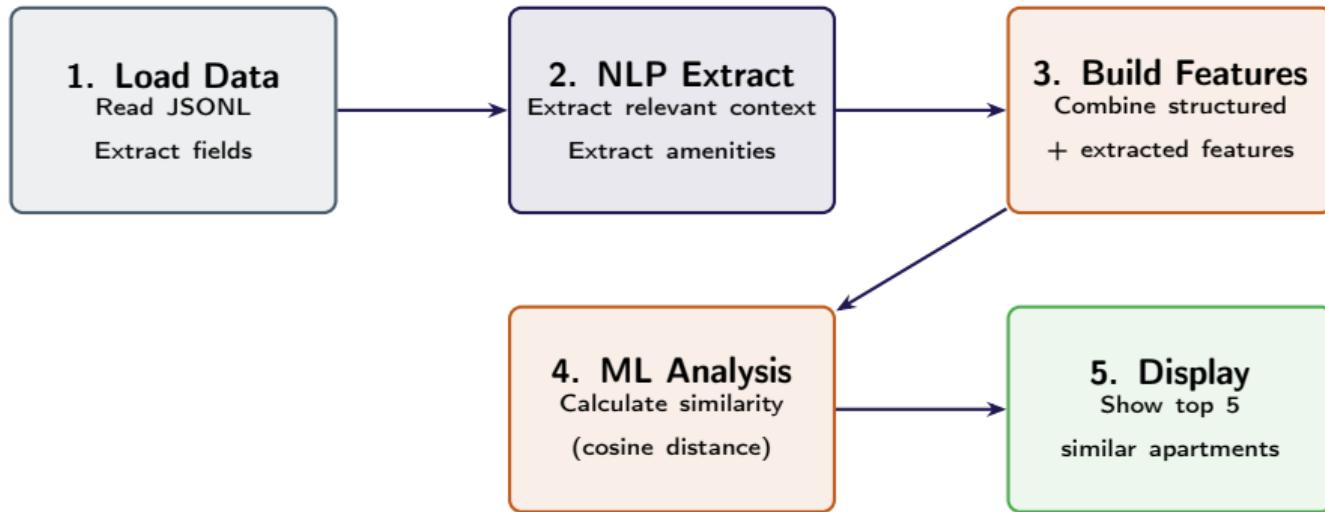
- **Incremental Integration**

Connect components one at a time, testing after each connection.

- *Example:* Data loading → Add NLP → Test → Add ML → Test → Add UI
- ✓ Always have a working version, know exactly what broke if something fails



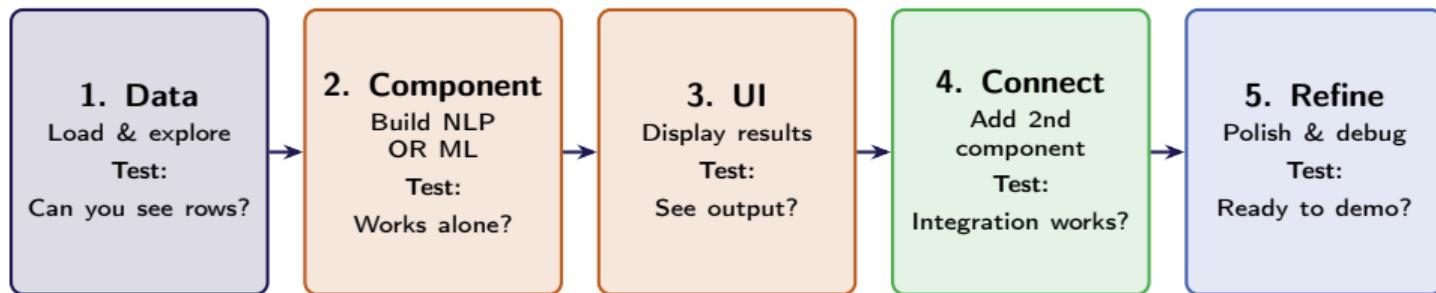
How components work together: A simple apartment recommendation system



Key Point: Each box is an independent function/module. You can:

- Replace step 2 (regex → LLM) without touching steps 4-5
- Swap step 4 (cosine → network distance) without changing steps 1-3
- Improve step 5 (table → interactive map) without touching analysis

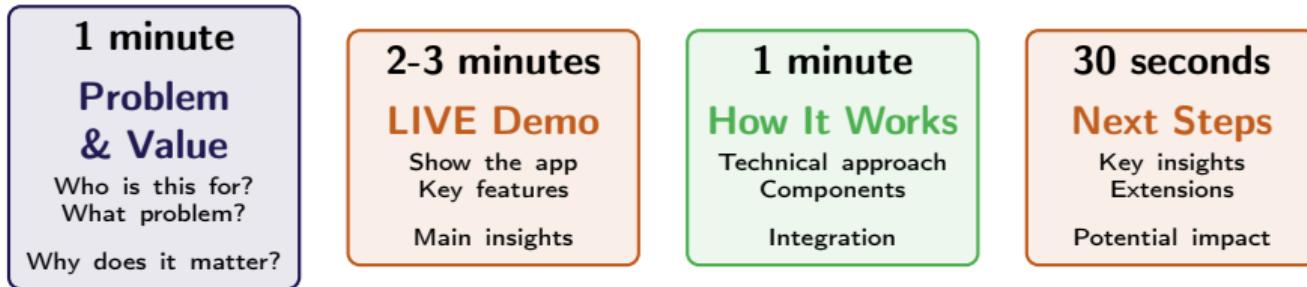
The key to success: Always have something working



Core Principle: Each step work independently and as intended. Test before moving forward.
Never integrate broken components.

Evaluation and Execution

Your pitch should tell a clear story: Problem → Solution → Impact



Do:

- Practice your demo beforehand
- Have backup screenshots
- Focus on business value
- Be concise and clear

Don't:

- Read code line-by-line
- Apologize for what's missing
- Go over time
- Debug live on stage



How your POC will be assessed:

Business Value - Clear problem, relevant solution, articulated value

Technical Execution - App works, appropriate methods, integrated

Presentation - Clear pitch, effective demo

Insight Quality - Novel findings

What you need to deliver, and how you can stand out:

Core Requirements (Everyone Must Deliver):

- ✓ **Text Analysis:** Extract features from descriptions (keywords, LLM extraction, embeddings)
- ✓ **ML/Analysis:** One working component (clustering, regression, recommendations, or network)
- ✓ **Application:** Working Streamlit app with interactivity
- ✓ **Value Proposition:** Clear problem statement and demo-ready pitch

Ways to Enhance (Choose What Fits Your Solution):

- **Deeper Integration:** Components genuinely enhance each other (e.g., network metrics as ML features)
- **Sophisticated Analysis:** Advanced NLP (semantic search), complex ML (ensemble models), or network analysis
- **Multiple Components:** Combine NLP + ML + Network in meaningful ways
- **Novel Insights:** Discover patterns or approaches beyond the obvious
- **Polished Experience:** Intuitive UI, compelling visualizations, professional presentation

Before you start building, remember these principles:

- **Start with Business Value, Not Technology**

Ask "What problem am I solving?" before "What algorithm should I use?"

- **Sketch Before Coding**

5 minutes planning architecture = 30 minutes saved debugging

- **Test Components Independently**

Don't wait until integration to see if things work

- **Use the Checkpoint to Course-Correct**

At 1 hour: Stop, evaluate, adjust if needed

- **Prepare Demo Backup**

Screenshots in case app breaks during presentation

- **Focus on Storytelling**

Why does this matter? Who benefits? What changes?